

Reducing Miss Rate

- Miss rate reduction represents the classical approach to improving cache behavior, therefore has received a lot of (possibly misplaced) interest over last decade.
- **Could you characterize the source of cache misses?**
- Consider the “Tree C’s” model,
 - Attempts to characterize misses into 3 classes,
 - Compulsory – cache at boot will always be empty.
 - First (cache) references are always misses (cold start).
 - Capacity – limited cache size
 - Entire program does not reside within cache.
 - Flushed blocks later require reloading.
 - Conflict – too many main memory (MM) blocks mapped to the same cache set (direct and set associative).
- Figure 1 qualifies miss rate on different cache sizes as a function of the three C’s for 1- (i.e. direct), 2-, 4- and 8-way set associative caches.
 - Compulsory misses – effectively contribute no impact and are independent of cache size;
 - Conflict misses – decreases with increasing associativity;
 - Capacity misses – decreases with increasing cache size;
 - Miss rate – decreases with increasing associativity.

Limitations of the 3 C’s model

- Only describes average cache miss behavior
 - Reason for individual misses is not available.
- No temporal information captured
 - Increasing associativity maximizes cache hits but may increase cache hit penalty!
- Following techniques for reducing miss rate need to be taken within the context of overall cache performance.

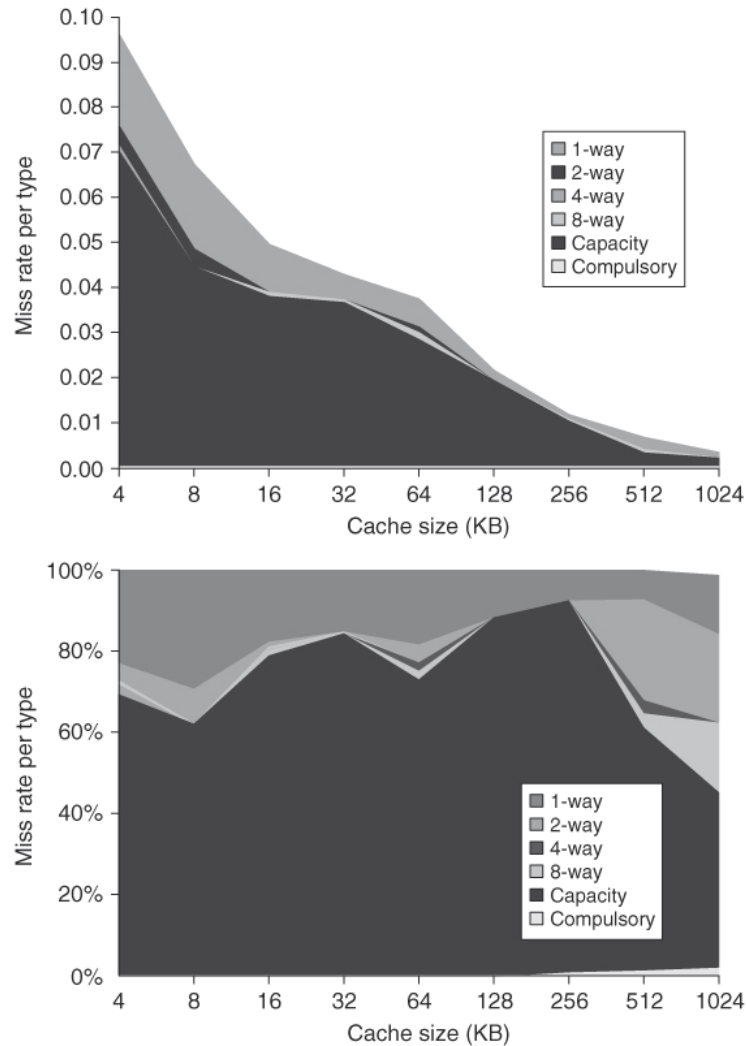


Figure 1: Total Miss Rate (top), Distribution of Miss Rates (bottom) as a function of the three C's

#1 – Larger Caches


- Motivation – let more data/ instructions reside in cache.
- Any disadvantage?
 - Longer hit time!!
- Table 1 lists average memory access time against block size for five cache sizes, with lowest time in bold.
- Pragmatics:
 - Use on off-chip caches, hit time already high and last defense against a call to disk

Table 1 – Average memory access time against block size for 5 cache sizes

Block Size	Miss Penalty	Cache Size			
		4K	16K	64K	256K
16	82	8.027	4.231	2.673	1.894
32	84	7.082	3.411	2.134	1.558
64	88	7.160	3.323	1.933	1.447
128	96	8.469	3.659	1.979	1.470
256	112	11.961	4.685	2.288	1.549

#2 – Larger Block Sizes

- Motivation – maximize spatial locality.
- Figure 2 plots miss rate as a function of block size on SPEC 92 benchmark with DECstation 5000.
 - Optimal block size appears to be a ratio of cache size.
- Advantages
 - Emphasize spatial locality of data/ instructions;
 - Reduction in compulsory misses
 - (from Figure 1 this is actually of minimum significance);
- Disadvantages
 - Increasingly sensitive to miss penalty as block size increases;
 - Conflict misses increase,
 - Reduced block diversity, more MM blocks mapped to same cache set (set associative).




© 2003 Elsevier Science (USA). All rights reserved.

Figure 2: Miss rate against block size for 5 different cache sizes.

#3 – Higher Associativity

- With respect to Figure 1,
 - 8-way set associative approximates fully associative;
 - 2:1 cache rule of thumb (for cache size < 128 K bytes)
 - miss rate of directly mapped cache of size N = miss rate of 2-way set associative cache of size $N / 2$.
- Comment
 - Access time increases with associativity, Figure 3.
 - Current practice is to utilize 1- to 4-way set associative caches.
 - Short CPU clock cycle rewards simple (1-way) cache designs – e.g. L1 cache
 - High miss penalties rewards higher associativity – e.g. L2 cache



© 2003 Elsevier Science (USA). All rights reserved.

Figure 3: Access time as a function of associativity and cache size.

#4 – Way Prediction

- As per dynamic branch prediction, each block in a set associative cache has a predictor.
 - Predictor selects which block of a set to try at the *next* cache access.
 - Tag comparison on predicted block conducted in parallel with the read.
 - If the prediction is incorrect,
 - CPU flushes the fetched instruction and the predictor is updated;
 - Continue to make tag comparisons for remaining blocks.

#5 – Compiler Optimizations

- Independent of specific hardware units.
- Concentrate on code re-organization for increasing the degree of temporal and spatial code locality.
- Consider two specific examples,
 - Loop interchange and Blocking.

(a) Loop Interchange


- Objective – Improve spatial locality
- Methodology
 - Cache does not have capacity for all matrix content.
 - Order of loop nesting provides good spatial locality if data accessed in the order in which it originally appears.
 - Cache thrashing results if cache capacity or conflicts stop all data appearing in cache “as one”.
- Consider,

Unoptimized code	Optimized code
<pre>for (j = 0; j < 100; j++) { for (i = 0; i < 5000; i++) { x[i][j] = 2 * x[i][j]; } }</pre>	<pre>for (i = 0; i < 5000; i++) { for (j = 0; j < 100; j++) { x[i][j] = 2 * x[i][j]; } }</pre>

(b) Blocking

- Objective – Improve Temporal locality of data.
- Problem
 - Interested in accessing same data in both row and column format.
 - Cache does not have the capacity for any more than a row or column.
 - i.e. reading entire row or column into cache is optimally wrong for one set of operations.
- Methodology
 - Complete all row and column operations on ‘block’ of data in one pass.
- Consider the following code,

```
for (i = 0; i < N; i++)
{
    for (j = 0; j < N; j++)
    {
        r = 0;
        for (k = 0; k < N; k++)
        {
            r = r + y[i][k] * z[k][j];
        }
        x[i][j] = r;
    }
}
```
- What is the linear math equivalent?
 - One value for ‘r’ requires all row of ‘y’ and entire column from ‘z’.
- Process summarized by figure 4,
- Potential for domination by capacity misses.
 - If the cache may hold all 3 N by N matrices, then no problem.
 - If cache may hold one N by N matrix (z) and one row (y) then flush each row of ‘y’ N times.
 - If less than $N + N \times N$ then likelihood of cache thrashing increases.
 - Worst case capacity misses of $2N^3 + N^2$ accesses for N^3 operations.




© 2003 Elsevier Science (USA). All rights reserved.

Figure 4: Naïve matrix multiplication (outer product) for index $i = 1$. (LHS is the result matrix)
Dark gray – recent access; Light gray – old access; White – no access.

- Solution

- Divide the cache into sub-blocks, size B by B;
- Value for result ‘r’ built up incrementally over all blocks.
- Inner loops compute over the sub-blocks, figure 5;



© 2003 Elsevier Science (USA). All rights reserved.

Figure 5: Block based outer product multiplication. Case of $B = 3$. (LHS is the result matrix)

- ‘B’ the blocking factor is selected to ensure that sub-blocks for all three variables (x, y, z) reside in cache (or registers).
- Number of capacity misses is now $2N^3 \div B + N^2$ accesses for N^3 operations.
 - I.e. speedup of ‘B’
 - Following code illustrates a possible ‘block based’ implementation.

```
for (jj = 0; jj < N; jj = jj + B)
{
    for (kk = 0; kk < N; kk = kk + B)
    {
        for (i = 0; i < N; i++)
        {
            for (j = jj; j < min(jj + B, N); j++)
            {
                r = 0;
                for (k = kk; k < min(kk + B, N); k++)
                {
                    r = r + y[i][k] * z[k][j];
                }
                x[i][j] = x[i][k] + r;
            }
        }
    }
}
```

Summary – Reducing Cache Miss Rate

- Three C's model,
 - Large block size – reduce compulsory misses;
 - Large cache size – reduce capacity misses;
 - High associativity – reduce conflict misses.
 - Model ignores effects on hit time and miss penalties.
- Way Prediction
 - Reduces hit time with higher set associativity designs.
- Compiler optimizations
 - Improve organization of memory accesses.
 - Require compile time knowledge of parameters.